

# ENSO Forecasting: MIFS Feature Selection and XGBoost Implementation

July 16, 2026

## IMPLEMENTATION

### 1. Mutual Information Feature Selection (MIFS)

MIFS is a greedy feature selection algorithm that balances relevance (mutual information with target) and redundancy (mutual information among selected features).

#### 1.1 Algorithm Overview

The MIFS score for feature  $X_{ij}$  is defined as:

$$S(X_{ij}) = I(X_{ij}; Y) - \frac{1}{|S|} \sum_{X_{kl} \in S} I(X_{ij}; X_{kl}) \quad (1)$$

where:

- $I(X_{ij}; Y)$  is the mutual information (relevance) between feature and target
- $I(X_{ij}; X_{kl})$  is the mutual information (redundancy) between feature pair
- $S$  is the current set of selected features

#### 1.2 Implementation Steps

##### Step 1: Compute Relevance

- For all  $\sim 20,000$  features in windowed dataset
- Use Kraskov k-NN estimator (sklearn's `mutual_info_regression`)
- k-neighbors parameter:  $k = 5$  for MI computation
- Training set only: 457 samples (window end  $\leq 2018-12$ )

**Step 2: Stratified Pre-filtering** Unlike flat top-N selection, stratified prefiltering uses per-variable-type quotas:

- Group features by variable type (e.g., `sst_anom_z`, `sohtc300_anom_z`)
- Per-type quota =  $\min(\text{cap\_max}, \max(\text{cap\_min}, \text{fraction} \times n_{\text{type}}))$
- Parameters: `fraction` = 0.4, `cap_min` = 15, `cap_max` = 30
- Result: Candidate pool of  $\leq 300$  features with all variable types represented
- Rationale: Prevents SST (strongest individual predictor) from monopolizing candidate pool; ensures mechanistically important but weak predictors (OHC, wind stress) are included

### Step 3: Greedy MIFS Loop

- Initialize:  $S = \emptyset$
- For each step  $t = 1, \dots, \min(50, |\text{candidates}|)$ :
  1. For each remaining candidate  $X_i$ :
  2. If  $|S| = 0$ :  $\text{score} = I(X_i; Y)$  (pure relevance)
  3. Else:  $\text{score} = I(X_i; Y) - \frac{1}{|S|} \sum_{X_j \in S} I(X_i; X_j)$  (redundancy penalty)
  4. Cache pairwise MI to avoid recomputation
- Select:  $X^* = \arg \max_i \text{score}(X_i)$
- Update:  $S \leftarrow S \cup \{X^*\}$
- Output: Ranked list of selected features in order

### Step 4: K-sweep with Validation

- For each  $K = 1, 2, \dots, 50$  in steps of 2:
- Take top- $K$  features from greedy ranking
- Train XGBoost (baseline: `n_est` = 200, `max_depth` = 3, `lr` = 0.1) on train set (457 samples)
- Evaluate on validation set (48 samples, window end 2019-01 to 2022-12)
- Compute RMSE on validation predictions
- Select:  $K^* = \arg \min_K \text{RMSE}_{\text{val}}(K)$




Figure 1: K-sweep validation: RMSE vs. number of features K. The optimal K minimizes validation error.

#### Step 5: Output Save:

- Full relevance ranking (all  $\sim 20,000$  features)
- MIFS-selected features in ranking order
- K-sweep results (RMSE for each K)
- Best K and final feature list

## 2. XGBoost Regression

XGBoost (eXtreme Gradient Boosting) is an ensemble tree model that sequentially adds trees to minimize loss.

### 2.1 Regularization-Focused Grid Search

To prevent overfitting observed in earlier runs, a **regularized grid** was designed:

| Parameter        | Baseline | Regularized     |
|------------------|----------|-----------------|
| max_depth        | 6        | <b>2, 3</b>     |
| n_estimators     | 500      | 100, 200, 300   |
| learning_rate    | 0.1      | 0.01, 0.05, 0.1 |
| subsample        | —        | 0.6, 0.7, 0.8   |
| colsample_bytree | —        | 0.6, 0.7, 0.8   |
| min_child_weight | —        | 5, 10, 20       |
| reg_alpha (L1)   | —        | 0, 0.5, 1.0     |
| reg_lambda (L2)  | —        | 1, 5, 10        |

Table 1: Regularized XGBoost grid (4,374 combinations)

## 2.2 Hyperparameter Tuning Process

### Stage 1: Grid Search on Validation Set

- Enumerate all parameter combinations ( $3 \times 3 \times 3 \times 3 \times 3 \times 3 \times 3 \times 3 = 6561$  when all expanded)
- Actually:  $2 \times 3 \times 3 \times 3 \times 3 \times 3 \times 3 \times 3 = 4374$  combos ( $\text{max\_depth} \in \{2, 3\}$ )
- For each combo: Train on full train set ( $457 + 48 = 505$  samples pooled), evaluate on validation set
- Metric: Validation RMSE
- Select best hyperparameters:  $\theta^* = \arg \min_{\theta} \text{RMSE}_{\text{val}}(\theta)$

### Stage 2: Final Evaluation on Test Set

- Re-train model with  $\theta^*$  on combined train+val set (505 samples)
- Evaluate on held-out test set (35 samples, window end 2023-01 to 2023-12+)
- Metrics: RMSE, ACC (anomaly correlation coefficient)

## 2.3 Anomaly Correlation Coefficient (ACC)

For ENSO forecasting, ACC is more operationally relevant than raw RMSE:

$$\text{ACC} = r(y_{\text{true}}, y_{\text{pred}}) = \frac{\text{Cov}(y_{\text{true}}, y_{\text{pred}})}{\sigma_{\text{true}} \cdot \sigma_{\text{pred}}} \quad (2)$$

- Pearson correlation between predicted and actual Niño 3.4 index
- Range:  $[-1, 1]$ ; higher is better
- Measures whether model gets the sign and phase correct

### Why ACC Matters for ENSO:

- Forecast skill depends on predicting El Niño/La Niña **phase** correctly
- A prediction  $0.05^\circ\text{C}$  off but correct sign  $>$  a prediction  $0.01^\circ\text{C}$  off but wrong sign

—

# RESULT AND ANALYSIS

## 1. MIFS Feature Selection Results

### 1.1 Stratified vs. Non-Stratified Prefiltering

| Model          | Prefilter Type | Best K | Top Features                      |
|----------------|----------------|--------|-----------------------------------|
| Baseline       | Flat top-300   | 22     | sossheig, t2m, sohtc300, so20chgt |
| Stratified     | Quota-based    | 10     | sossheig, t2m, sst, msl           |
| Tuned          | Tuned (flat)   | 18     | sossheig, sohtc300, t2m, so20chgt |
| Regularized    | Tuned (flat)   | 18     | (same as Tuned)                   |
| Stratified+Reg | Quota-based    | 10     | sossheig, t2m, sst, msl, avg_inss |

Table 2: MIFS feature counts and top predictors across models

#### Key Observations:

- **Flat prefiltering** selects 18-30 features; typically dominated by strongest individual predictors (SST, OHC)
- **Stratified prefiltering** selects fewer (10) but guarantees per-type representation
- **Top-ranked feature across all runs:** `sossheig_anom_z_(0.0,192.0)_lag0` (sea surface height at Niño 3.4 box, current month)
- **Consistent predictors:** Ocean heat content (sohtc300), thermocline (so20chgt), atmospheric (t2m, msl)

### 1.2 Redundancy Penalty Effect

Why does MIFS select so few features from  $\sim 20,000$ ?

- **Spatial correlation:** All grid cells in same ocean basin are highly correlated
- Once one SST cell is selected, all nearby cells get large redundancy penalty
- MIFS greedy loop effectively clusters spatially-correlated features and picks one representative
- With only 457 training samples, adding more features causes validation RMSE to increase (curse of dimensionality)
- K-sweep finds sweet spot: typically 10-22 features minimize validation error

—

## 2. XGBoost Performance Across Models

### 2.1 Baseline vs. Tuned vs. Regularized

| Model          | Features | Train RMSE        | Val RMSE | Test RMSE | Test ACC        |
|----------------|----------|-------------------|----------|-----------|-----------------|
| Baseline       | 22       | not logged        | 0.6821   | 0.8554    | 0.5608          |
| Stratified     | 10       | not logged        | 0.7089   | 0.8949    | 0.4972          |
| Tuned          | 18       | <b>0.0005</b> (!) | 0.6054   | 0.8906    | 0.8078          |
| Regularized    | 18       | 0.3934 ✓          | 0.6000   | 0.9077    | <b>0.8163</b> ✓ |
| Stratified+Reg | 10       | 0.3581 ✓          | 0.6534   | 0.9466    | 0.5872          |

Table 3: Model performance comparison (Train/Val/Test RMSE and ACC)

### 2.2 Key Results

#### Baseline (Flat MIFS + Standard XGB):

- 22 selected features
- Test RMSE: 0.8554°C (predictions off by 0.86°C on average)
- Test ACC: 0.5608 (moderate skill at predicting sign)
- No train/val data logged; unclear generalization

#### Tuned (Tuned MIFS + Expanded XGB Grid):

- 18 selected features
- **Severe overfitting:**
  - Train RMSE = 0.0005 (essentially perfect on training data)
  - Train ACC = 1.00 (correlates perfectly with training targets)
  - Test RMSE = 0.8906 (generalizes poorly)
  - Test ACC = 0.8078 (good on test despite overfitting)
- Train-test gap: +0.8901 RMSE (model memorized training set)
- **Not recommended for operational use**




Figure 2: Tuned model test predictions: actual vs. predicted Niño 3.4 index (2023-2025).




Figure 3: Tuned model feature importances: shows which features XGBoost relied on most.

### Regularized (Tuned MIFS + Regularized XGB Grid):

- 18 selected features (same as Tuned, using reused MIFS from Tuned)
- **Controlled overfitting:**

- Train RMSE = 0.3934 (learning, not memorizing)
- Train ACC = 0.9139 (strong but not perfect)
- Test RMSE = 0.9077 (only +0.517 worse than train)
- Test ACC = **0.8163** (best among all models)
- Train-test gap: +0.5143 RMSE (reasonable, indicates good generalization)
- Regularization parameters:
  - max\_depth = 2 (shallow trees only)
  - min\_child\_weight = 20 (high leaf threshold)
  - reg\_lambda = 5 (strong L2 penalty)
  - subsample = 0.6, colsample\_bytree = 0.6 (aggressive subsampling)
- **Recommended for operational deployment**




Figure 4: Regularized model overfitting analysis: train vs. validation performance across epochs.






Figure 5: Regularized model test predictions: actual vs. predicted Niño 3.4 index (2023-2025). Higher correlation than tuned model.




Figure 6: Regularized model feature importances: top predictors from stratified MIFS selection.

### Stratified+Regularized (Quota-based MIFS + Regularized XGB Grid):

- 10 selected features (smallest model)
- Combines stratified quota-based MIFS with anti-overfitting regularization

- **Performance metrics:**
  - Train RMSE = 0.3581 (controlled learning)
  - Train ACC = 0.9297 (strong)
  - Val RMSE = 0.6534
  - Val ACC = 0.4038 (lower than expected)
  - Test RMSE = 0.9466 (highest among all models)
  - Test ACC = 0.5872 (moderate)
- Train-test gap: +0.5885 RMSE (good generalization)
- Best hyperparameters:
  - max\_depth = 2 (very shallow)
  - min\_child\_weight = 20 (high leaf threshold)
  - reg\_lambda = 5 (strong L2 penalty)
  - subsample = 0.6, colsample\_bytree = 0.8
- **Trade-off:** Most concise model (only 10 features) with good generalization, but lower test ACC than Regularized (18 features)

## 2.3 Model Comparison Summary

| Metric              | Winner                           |
|---------------------|----------------------------------|
| Best Test RMSE      | Baseline (0.8554)                |
| Best Test ACC       | Regularized (0.8163)             |
| Fewest Features     | Stratified+Reg & Stratified (10) |
| Best for Deployment | <b>Regularized</b>               |

Table 4: Recommended model: Regularized (best balance of RMSE, ACC, and robustness)

## 3. Analysis

### 3.1 Why Regularized Outperforms Tuned

**The Tuned model problem:**

$$\text{Train RMSE} = 0.0005 \Rightarrow \text{Model has memorized every training sample} \quad (3)$$

When training data is only 457 samples and model has  $\sim 100$  features  $\times$  6 levels (boosting depth), the model has enough capacity to fit every point perfectly. This is useless for forecasting new months.

**The Regularized model solution:**

- **Shallow trees (max\_depth=2):** Cannot fit complex, high-dimensional patterns
- **High min\_child\_weight (20):** Prevents leaf nodes for single samples
- **L2 penalty (reg\_lambda=5):** Shrinks leaf weight magnitudes
- **Row/column subsampling:** Each tree sees only 60% of rows/features
- Train RMSE stays  $\approx 0.39$  (real learning, not memorization)
- Test performance degrades gracefully to 0.91 RMSE (plausible generalization)

### 3.2 Anomaly Correlation vs. RMSE

For operational ENSO forecasting, **\*\*ACC often matters more than raw RMSE\*\***:

- **Baseline:** Lower RMSE (0.8554) but moderate ACC (0.5608)
  - Can predict magnitude but not always the phase (El Niño vs. La Niña)
- **Regularized:** Slightly higher RMSE (0.9077) but highest ACC (0.8163)
  - Gets the sign/phase right in 81.6% of predictions
  - More useful for seasonal forecasting

#### Example:

- True Niño 3.4 = +1.5C (El Niño)
- Baseline predicts: +1.48C (RMSE = 0.02, but if internally correlated poorly, ACC low)
- Regularized predicts: +1.42C (RMSE = 0.08, but correlation with true is strong, ACC high)
- For forecasting, Regularized is more reliable

### 3.3 Stratified vs. Flat MIFS Trade-off

- Flat (Baseline, Tuned, Regularized): Selects 18-22 features, mostly SST + OHC
- Stratified (Stratified, Stratified+Reg): Selects 10 features, guarantees all variable types
  - More concise model (fewer parameters)
  - More physically interpretable (all known ENSO predictors present)
  - Includes weaker-but-important predictors: avg\_inss (north-south wind stress)
- **Performance comparison:**
  - Stratified (no regularization): Test RMSE=0.8949, ACC=0.4972
  - Stratified+Reg (with regularization): Test RMSE=0.9466, ACC=0.5872

- Regularized (flat MIFS, 18 features): Test RMSE=0.9077, ACC=0.8163
- **Key finding:** Stratified MIFS with only 10 features achieves moderate performance (ACC=0.59) but falls short of flat MIFS + regularization (ACC=0.82)
- **Interpretation:** The highly correlated SST grid cells selected by flat MIFS carry more predictive signal than the quota-enforced diversity

### 3.4 Feature Count vs. Performance

- **10 features (Stratified+Regularized):** ACC=0.59 — too few, missing key patterns
- **18 features (Regularized):** ACC=0.82 — sweet spot for 457 training samples
- **22 features (Baseline):** ACC=0.56 — without regularization, starts overfitting